

NEURAL NETWORK AND DEEP LEARNING ASSIGNMENT-7

GITHUB LINK: <https://github.com/aknomula/DL-NN-Assign7.git>

RECORDINGLINK: https://drive.google.com/file/d/1PKoyuGA6Y21GqHs0_dZ-Ei2CmrolMiNj/view?usp=drive_link

Use Case Description:

Image Classification with CNN

1. Training the model
2. Evaluating the model

Programming elements:

1. About CNN
2. Hyperparameters of CNN
3. Image classification with CNN

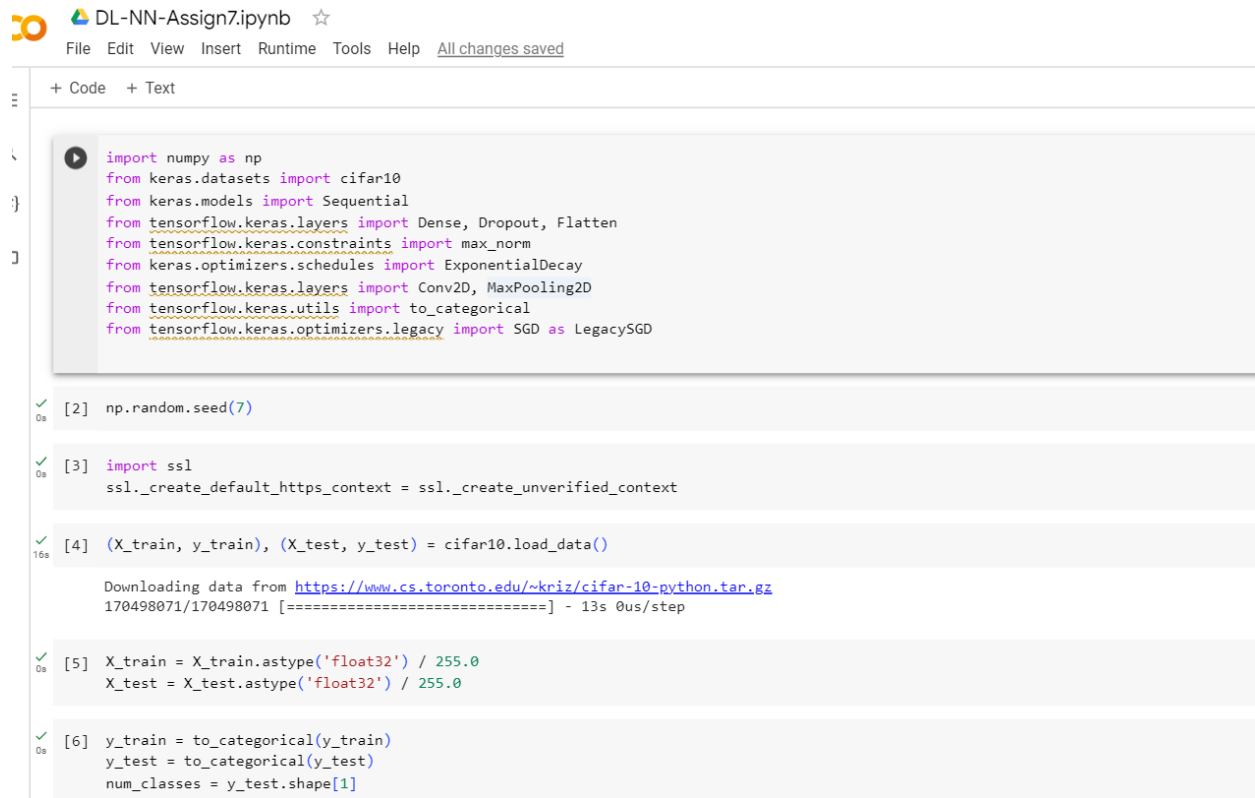
In class programming:

1. Follow the instruction below and then report how the performance changed.(apply all at once)

- Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2 .
- Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2 .
- Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.

- Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2.
- Flatten layer.
- Dropout layer at 20%.
- Fully connected layer with 1024 units and a rectifier activation function.
- Dropout layer at 20%.
- Fully connected layer with 512 units and a rectifier activation function.
- Dropout layer at 20%.
- Fully connected output layer with 10 units and a Softmax activation function

Did the performance change?



The screenshot shows a Jupyter Notebook titled "DL-NN-Assign7.ipynb". The interface includes a menu bar with options: File, Edit, View, Insert, Runtime, Tools, Help, and a link for "All changes saved". Below the menu bar, there are tabs for "+ Code" and "+ Text". The notebook contains several code cells:

- Cell 1:** Imports necessary libraries: `import numpy as np`, `from keras.datasets import cifar10`, `from keras.models import Sequential`, `from tensorflow.keras.layers import Dense, Dropout, Flatten`, `from tensorflow.keras.constraints import max_norm`, `from keras.optimizers.schedules import ExponentialDecay`, `from tensorflow.keras.layers import Conv2D, MaxPooling2D`, `from tensorflow.keras.utils import to_categorical`, and `from tensorflow.keras.optimizers import SGD as LegacySGD`.
- Cell 2:** `np.random.seed(7)`
- Cell 3:** `import ssl` and `ssl._create_default_https_context = ssl._create_unverified_context`
- Cell 4:** `(X_train, y_train), (X_test, y_test) = cifar10.load_data()`. The output shows the download progress for the CIFAR-10 dataset from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>.
- Cell 5:** `X_train = X_train.astype('float32') / 255.0` and `X_test = X_test.astype('float32') / 255.0`
- Cell 6:** `y_train = to_categorical(y_train)`, `y_test = to_categorical(y_test)`, and `num_classes = y_test.shape[1]`

```

model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=max_norm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

```

```
[8] #sgd = SGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
```

```

sgd = LegacySGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

```

=====
Total params: 4210090 (16.06 MB)
Trainable params: 4210090 (16.06 MB)
Non-trainable params: 0 (0.00 Byte)

```

None

```

[9] epochs = 5
batch_size = 32
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)

```

```

Epoch 1/5
1563/1563 [=====] - 19s 7ms/step - loss: 1.7193 - accuracy: 0.3756 - val_loss: 1.3652 - val_accuracy: 0.5145
Epoch 2/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.3508 - accuracy: 0.5155 - val_loss: 1.2663 - val_accuracy: 0.5463
Epoch 3/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.1818 - accuracy: 0.5800 - val_loss: 1.1044 - val_accuracy: 0.6070
Epoch 4/5
1563/1563 [=====] - 10s 6ms/step - loss: 1.0602 - accuracy: 0.6252 - val_loss: 1.0444 - val_accuracy: 0.6300
Epoch 5/5
1563/1563 [=====] - 10s 6ms/step - loss: 0.9510 - accuracy: 0.6647 - val_loss: 1.0254 - val_accuracy: 0.6429
<keras.src.callbacks.History at 0x7d7f1190550>

```

```

[10] scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

```

Accuracy: 64.29%

```

import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.constraints import max_norm
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import SGD as LegacySGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=max_norm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=max_norm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

```

```

# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = LegacySGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
#sgd = LegacySGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

```

```

Model: "sequential_3"
Layer (type) Output Shape Param #
-----
conv2d_14 (Conv2D) (None, 32, 32, 32) 896
dropout_14 (Dropout) (None, 32, 32, 32) 0
conv2d_15 (Conv2D) (None, 32, 32, 32) 9248
max_pooling2d_7 (MaxPooling2D) (None, 16, 16, 32) 0
conv2d_16 (Conv2D) (None, 16, 16, 64) 18496
dropout_15 (Dropout) (None, 16, 16, 64) 0
conv2d_17 (Conv2D) (None, 16, 16, 64) 36928
max_pooling2d_8 (MaxPooling2D) (None, 8, 8, 64) 0
conv2d_18 (Conv2D) (None, 8, 8, 128) 73856
dropout_16 (Dropout) (None, 8, 8, 128) 0
conv2d_19 (Conv2D) (None, 8, 8, 128) 147584
max_pooling2d_9 (MaxPooling2D) (None, 4, 4, 128) 0
flatten_3 (Flatten) (None, 2048) 0
dropout_17 (Dropout) (None, 2048) 0
dense_8 (Dense) (None, 1024) 2098176
dropout_18 (Dropout) (None, 1024) 0
dense_9 (Dense) (None, 512) 524800
dropout_19 (Dropout) (None, 512) 0
dense_10 (Dense) (None, 10) 5130
-----
Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)
-----
/usr/local/lib/python3.10/dist-packages/keras/src/optimizers/legacy/gradient_descent.py:114: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
None
Epoch 1/5
1563/1563 [=====] - 16s 9ms/step - loss: 1.9086 - accuracy: 0.2964 - val_loss: 1.5709 - val_accuracy: 0.4328
Epoch 2/5
1563/1563 [=====] - 13s 8ms/step - loss: 1.5088 - accuracy: 0.4494 - val_loss: 1.4358 - val_accuracy: 0.4813
Epoch 3/5

```

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.

```

[] import numpy
# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
# Convert the predictions to class labels
predicted_labels = numpy.argmax(predictions, axis=1)
# Convert the actual labels to class labels
actual_labels = numpy.argmax(y_test[:4], axis=1)

# Print the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels: ", actual_labels)

1/1 [=====] - 0s 20ms/step
Predicted labels: [3 8 8 0]
Actual labels:   [3 8 8 0]

```

3. Visualize Loss and Accuracy using the history object

```
import matplotlib.pyplot as plt

# Plot the training and validation loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

# Plot the training and validation accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='lower right')
plt.show()
```

